

On the Scalability of Ad Hoc Routing Protocols

César A. Santiváñez
Internet. Research Dept.
BBN Technologies
Cambridge, MA
csantiva@bbn.com

Bruce McDonald
Elec. & Comp. Eng. Dept.
Northeastern University
Boston, MA
mcdonald@ece.neu.edu

Ioannis Stavrakakis
Dept. of Informatics
University of Athens
Athens, Greece
istavrak@di.uoa.gr

Ram Ramanathan
Internet. Research Dept.
BBN Technologies
Cambridge, MA
ramanath@bbn.com

Abstract—A novel framework is presented for the study of scalability in ad hoc networks. Using this framework, the first asymptotic analysis is provided with respect to network size, mobility, and traffic for each fundamental class of ad hoc routing algorithms. Protocols studied include the following: Plain Flooding (PF), Standard Link State (SLS), Dynamic Source Routing (DSR), Hierarchical Link State (HierLS), Zone Routing Protocol (ZRP), and Hazy Sighted Link State (HSLS). It is shown that PF and ZRP scale better with mobility, SLS and ZRP scale better with respect to traffic, and HSLS scales better with respect to network size. The analysis provides deeper understanding of the limits and trade-offs inherent in mobile ad hoc network routing. Our analysis is complemented with a simulation experiment comparing HSLS and HierLS. An important contribution of this paper is that HSLS is an scalable, easy-to-implement, alternative to hierarchical approaches for large ad hoc networks.

I. INTRODUCTION

Routing protocols for ad hoc networks have been the subject of extensive research over the past several years. Recently, practical applications such as intelligent sensor networks have focused attention on understanding the issues and tradeoffs in network scalability. An important question that arises is : *which routing protocol scales the best?* The typical answer is: *it depends*. Unfortunately, the networking community lacks a tenet for understanding the fundamental properties and limitations of ad hoc networks. Hence, a fundamental understanding of *what* scalability depends on, and *how* is currently lacking.

One reason for this shortcoming is a lack of sufficient research aimed at general principles and analytical modeling. Scalability and other performance aspects of ad hoc routing have been studied predominantly via simulations (e.g. [1], [2], [3]), versus theoretical analyses. Simulation results, although extremely useful, are often limited in scope to specific scenarios. Thus, they often fail to produce results that provide the depth of understanding of the limitations of the protocols and their dependence on system parameters and environmental factors desired by researchers. The lack of much needed theoretical analysis in this area is due, we believe, in part to the lack of a common platform to base theoretical comparisons on, and in part due to the abstruse nature of the problem.

This paper focuses on the development of principles and methodologies for the analysis and design of scalable routing strategies for ad hoc networks. Analytical models are developed and results are presented that provide significant insight into the aforementioned dependency and the general performance characteristics of the most important classes of ad hoc network routing algorithms. The theoretical models developed establish the basis for an unbiased analysis and comparison of the relative scalability of several proposed routing protocols.

The first precise (asymptotic) expressions reflecting the im-

pact of network size, traffic intensity and mobility on protocol performance are developed in this paper. Analytical results are presented for a representative set of state-of-the-art protocols in the literature, including : no routing – Plain Flooding (PF), proactive – Standard Link State (SLS), reactive – Dynamic Source Routing (DSR) [7], hybrid – Zone Routing Protocol (ZRP) [9], hierarchical – Hierarchical Link State (HierLS) [8], and limited dissemination – Hazy Sighted Link State (HSLS) [11], techniques. As such, the results provide researchers with improved understanding of the limits and trade-offs inherent in ad hoc network routing. A significant result is that, under the assumptions of this work, HSLS—while being easier to implement — scales better than HierLS and ZRP with respect to network size. This analytical result is validated with simulation analysis comparing HSLS and HierLS. Thus, another important contribution of this work is to show that HSLS is an scalable, more efficient alternative than hierarchical approaches for routing in large ad hoc networks.¹

Despite limited prior related theoretical work, there have been notable exceptions. In [4] analytical and simulation results are integrated in a study that provides valuable insight into comparative protocol performance. However, it fails to deliver a final analytical result, deferring instead to simulation. Thus, it is difficult to fully understand the interactions among system parameters. The present work closes this gap and provides an understanding of the dynamic interaction among network parameters.

The asymptotic capacity of a fixed wireless network was studied in [5]; however, it did not include routing overhead. In contrast, we focus on *total overhead* (defined later), which includes routing overhead. The impact of mobility on network capacity was studied in [6]. They showed that given no restriction on memory size and arbitrarily long delays, mobility increases network capacity. This research, however, focuses on practical scenarios, wherein, delay cannot grow arbitrarily large and mobility reduces the network capacity (degrading performance).

The remainder of this paper is organized as follows: In Section-II we characterize the (*total overhead*) metric and the network model used. Sections-III - VIII present analysis of the asymptotic performance of PF, SLS, DSR, HierLS, ZRP, and HSLS respectively. Comparison of protocol performance is discussed in Section-IX, focusing on HierLS and HSLS under large network size and including simulation results. Finally, conclusions are presented in Section-X

¹However, it should be noted that HSLS requires more memory than HierLS. As network size increases HSLS's memory requirements may become the limiting factor.

II. MODELING PRELIMINARIES

This section presents the model assumptions and definitions employed in our analysis.

A. Network model

The following notation will be utilized in this paper: Let N be the number of nodes in the network, d be the average in-degree, L be the average path length (in hops) over all source destination pairs, λ_{lc} be the expected number of link status changes that a node detects per second, λ_t be the average traffic rate that a node generates in a second (in bps), λ_s be the average number of new sessions generated by a node per second, and $data$ be the average data packet size (in bits). This work uses the same set of assumptions, based on geographical reasoning, that were presented and discussed in [11], [12], [14], [15], which are reproduced below for the sake of clarity.²

- a.1. As the network size increases, the average in-degree d remains constant.
- a.2. Let A be the area covered by the N nodes of the network, and $\sigma = N/A$ be the network average density. Then, the expected (average) number of nodes inside an area A_1 is approximately σA_1 .
- a.3. The number of nodes that are at distance of k or less hops away from a source node increases (on average) as $\Theta(d k^2)$. The number of nodes exactly k hops away increases as $\Theta(d k)$.
- a.4. The maximum and average paths (in hops) among nodes in a connected subset of n nodes both increase as $\Theta(\sqrt{n})$. In particular, the maximum path length across the entire network and the average path length across the network (L) increase as $\Theta(\sqrt{N})$.
- a.5. The traffic that a node generates each second (λ_t and λ_s), is independent of the network size N (number of destinations). As the network size increases, the total amount of data transmitted/received by a single node remains constant, but the number of destinations increases (traffic diversity will increase).
- a.6. For a given source node, all possible destinations ($N - 1$ nodes) are equiprobable. The traffic from one node to a given destination decreases as $\Theta(1/N)$.
- a.7. Link status changes are due to mobility. λ_{lc} is directly proportional to the relative node speed.
- a.8. Mobility models : time scaling. Let $g_{0/1}(x, y)$ be the probability distribution function of a node position at time 0 second, given that it is known that the node position at time 1 will be $(0, 0)$. Then, the probability distribution function of a node position at time $t < t_1$ given that the node will be at the position (x_{t_1}, y_{t_1}) at time t_1 , is given by $g_{t/t_1}(x, y, x_{t_1}, y_{t_1}) = \frac{1}{(t-t_1)^2} g_{0/1}\left(\frac{x-x_{t_1}}{t_1-t}, \frac{y-y_{t_1}}{t_1-t}\right)$.

Assumptions a.1 - a.8 represents a well-defined network model, still general enough to include most of the typical networking scenarios. The reader is referred to [11], and [12] for a discussion on these assumptions.

It should be noted that in the case any of the above assumptions does not hold for a particular class of networks, alternative

²Standard asymptotic notation is employed. A function $f(n) = \Omega(g(n))$ [similarly, $f(n) = O(g(n))$] if there exists constants c_1 and n_1 [similarly, c_2 and n_2] such that $c_1 g(n) \leq f(n)$ [similarly $f(n) \leq c_2 g(n)$] for all $n \geq n_1$ [similarly, $n \geq n_2$]. Also, $f(n) = \Theta(g(n))$ if and only if $f(n) = \Omega(g(n))$ and $f(n) = O(g(n))$.

expressions may be derived by following the same methodology set forth in this paper.

B. Definitions: Total Overhead and Scalability

B.1 Total Overhead

Traditionally, the term *overhead* has been used in relation to the *control overhead*, that is, the amount of bandwidth required to construct and maintain a route. However, as shown in [11], a protocol's *control overhead* alone is not sufficient for assessing system performance, as it fails to account for the impact of sub-optimal routes. What is needed is a single metric that is able to capture the routing protocol impact on network performance. For bandwidth-constrained systems, the *total overhead* introduced in [11] and discussed below represents such a metric.

First, the *minimum traffic load* of the network must be defined, as follows:

Definition 1: The minimum traffic load of a network, is the minimum amount of bandwidth required to forward packets over the shortest distance (in number of hops) paths available, assuming all the nodes have instantaneous a priori full topology information.

The above definition is independent of the routing protocol being employed, since it does not include the control overhead but assumes that all the nodes are provided *a priori* global information. It should be noted that it is possible that in fixed networks a node is provided with static optimal routes, and therefore there is no bandwidth consumption above the *minimum traffic load*. On the other hand, in mobile scenarios this is hardly possible. Due to the unpredictability of the movement patterns and the topology they induce, even if static routes are provided so that no control packets are needed, it is extremely unlikely that the static routes so forced remain being the optimal ones during the entire network lifetime. Thus, since sub-optimal routes are present, the actual network bandwidth usage would be greater than the *minimum traffic load* value. This motivates the following definition of a routing protocol *total overhead*.

Definition 2: The total overhead induced by a routing protocol is the difference between the total amount of bandwidth actually consumed by the network running such routing protocol minus the minimum traffic load that would have been required should the nodes had a priori full topology information.

Thus, the actual bandwidth consumption in a network will be the sum of a protocol independent term, the *minimum traffic load*, and a protocol dependent one, the *total overhead*. Effective routing protocols should try to reduce the second term (*total overhead*) as much as possible.

The different sources of overhead that contribute to the *total overhead* may be grouped and expressed in terms of *reactive*, *proactive*, and *sub-optimal routing overheads*. All of these sources of overhead has been considered in the past, but the *total overhead* represents the first metric that successfully combines all of them in a unified framework, allowing a tractable model to be derived.

The *reactive overhead* of a protocol is the amount of bandwidth consumed by the specific protocol to build paths from a source to a destination, *after* a traffic flow to that destination has been generated at the source. In static networks, the reactive

overhead is a function of the rate of generation of new flows. In dynamic (mobile) networks, however, paths are (re)built not only due to new flows but also due to link failures in an already active path. Thus, in general, the reactive overhead is a function of both traffic and topology change.

The *proactive overhead* of a protocol is the amount of bandwidth consumed by the protocol in order to propagate route information *before* it is needed. This may take place periodically and/or in response to topological changes.

The *sub-optimal routing overhead* of a protocol is the difference between the bandwidth consumed when transmitting data from all the sources to their destinations using the routes determined by the specific protocol, and the bandwidth that would have been consumed should the data have followed the shortest available path(s). For example, consider a source that is 3 hops away from its destination. If a protocol chooses to deliver one packet following a k ($k > 3$) hop path (maybe because of out-of-date information), then $(k-3)*\text{packet_length}$ bits will need to be added to the sub-optimal routing overhead.

The *total overhead* provides an unbiased metric for performance comparison that reflects bandwidth consumption. Despite increasing efficiency at the physical and MAC-layers, bandwidth is likely to remain a limiting factor in terms of scalability, which is a crucial element for successful implementation and deployment of ad hoc networks. The authors recognize that *total overhead* may not fully characterize all the performance aspects relevant to specific applications. However, it can be used without loss of generality as it is proportional to factors including energy consumption, memory and processing requirements, and, furthermore, delay constraints have been shown to be expressed in terms of an equivalent bandwidth [13].

B.2 Scalability

This work is aimed at the study of the scalability properties of routing protocols for ad hoc networks. However, currently there is not a clear definition of scalability. Indeed, scalability has a different meaning for different people. Thus, we need to define the exact meaning of this term.

Definition 3: Scalability is the ability of a network to support the increase of its limiting parameters.³

Thus, scalability is a property. In order to quantify this property, we use the concept of *minimum traffic load* (definition 1) to define the *network scalability factor* as follows:

Definition 4: Let $\text{Tr}(\lambda_1, \lambda_2, \dots)$ be the minimum traffic load experienced by a network under parameters $\lambda_1, \lambda_2, \dots$ (e.g. network size, mobility rate, data generation rate, etc.). Then, the network scalability factor of such a network, with respect to a parameter λ_i (Ψ_{λ_i}) is defined to be :

$$\Psi_{\lambda_i} \stackrel{\text{def}}{=} \lim_{\lambda_i \rightarrow \infty} \frac{\log \text{Tr}(\lambda_1, \lambda_2, \dots)}{\log \lambda_i}$$

The *network scalability factor* is a number that asymptotically relates the increase in network load to the different network parameters. For the class of mobile ad hoc networks under study

³The *limiting parameters* of a network are those parameter – as for example mobility rate, traffic rate, and network size, etc. – whose increase causes the network performance to degrade. On the remainder of this work only limiting parameters will be considered, and therefore the term ‘parameter’ will be used in lieu of the term ‘limiting parameter’.

(assumptions a.1 - a.8), the *minimum traffic load* $\text{Tr}(\lambda_{lc}, \lambda_t, N)$ is $\Theta(\lambda_t N^{1.5})$, ⁴ and therefore $\Psi_{\lambda_{lc}} = 0$, $\Psi_{\lambda_t} = 1$, and $\Psi_N = 1.5$.

The *network scalability factor* may be used to compare the scalability properties of different networks (wireline, mobile ad hoc, etc.), and as a result of such comparisons we can say that one class of networks scales better than the other. However, if our desire is to assess whether a network is *scalable* (an adjective) with respect to a parameter λ_i , then the *network rate* dependency on such a parameter must be considered.

Definition 5: The network rate R^{net} of a network is the maximum number of bits that can be simultaneously transmitted in a unit of time.

For the *network rate* (R^{net}) computation all successful link layer transmissions must be counted, regardless of whether the link layer recipient is the final network-layer destination or not.

Definition 6: A network is said to be scalable with respect to the parameter λ_i if and only if, as the parameter λ_i increases, the network’s minimum traffic load does not increase faster than the network rate (R^{net}) can support. That is, if and only if:

$$\Psi_{\lambda_i} \leq \lim_{\lambda_i \rightarrow \infty} \frac{\log R^{\text{net}}(\lambda_1, \lambda_2, \dots)}{\log \lambda_i}$$

For example, it has been proved that in mobile ad hoc networks $\Theta(N)$ successful transmissions can be scheduled simultaneously (see for example [5], [6]). The class of networks under study in this work (i.e. resulting from applying power control techniques) are precisely the class of networks that achieves that maximum *network rate*. Thus, in order for mobile ad hoc network to be regarded as scalable with respect to network size, we will need $\Psi_N \leq 1$. Unfortunately this is not the case, and as a consequence ad hoc networks under assumption a.1 through a.8 are not *scalable* with respect to network size⁵. Wireline networks, in the other hand, if fully connected may have $\Psi_N = 1$, and therefore they are potentially scalable (in the bandwidth sense defined here) with respect to network size. Note however, that this scalability requires the nodes’ degree to grow without bound, which may be prohibitively expensive.

Similarly, since the *network rate* does not increase with mobility or traffic load, then a network will be scalable w.r.t. mobility and traffic if and only if $\Psi_{\lambda_{lc}} = 0$ and $\Psi_{\lambda_t} = 0$, respectively. Thus, the networks under this study are *scalable* w.r.t. mobility, but are not *scalable* w.r.t. traffic.

Note that similar conclusions may be drawn for scalability w.r.t. additional parameters as for example network density, transmission range ℓ , etc. that are not being considered in our analysis. For example, as transmission range increases (and assuming a infinite size network with regular density) the spatial

⁴Each node generate λ_t bits per seconds, that must be retransmitted (in average) L times (hops). Thus, each node induce a load of $\lambda_t L$, which after adding all the nodes results in a $\text{Tr}(\lambda_{lc}, \lambda_t, N) = \lambda_t N L$. Since, by assumption a.4 L is $\Theta(\sqrt{N})$, the above expression is obtained.

⁵It has been shown in [6] that if the network applications can support infinitely long delays and the mobility pattern is completely random, then the average path length may be reduced to 2 ($\Theta(1)$) regardless of network size and, as a consequence, that network *scalability factor* with respect to network size Ψ_N is equal to 1. Thus, those ad hoc networks (random mobility and capable of accepting infinitely long delays) are the only class of ad hoc networks that are scalable with respect to network size. This work does not consider that class of networks since they have no practical relevance.

reuse decreases and as a consequence *network rate* decreases as rapidly as ℓ^2 . Thus, Ψ_ℓ should be lower than -2 for the network to be deemed *scalable*. Since the *minimum traffic load* will only decrease linearly w.r.t. ℓ (paths are shortening), $\Psi_\ell = -1$, and therefore ad hoc networks are not scalable w.r.t. transmission range.⁶

Now, after noticing that mobile ad hoc networks are not *scalable* with respect to size and traffic, one may ask the meaning of regarding a routing protocol *scalable*. The remaining of this subsection will clarify this meaning.

Definition 7: Routing protocol's scalability is the ability of a routing protocol to support the continuous increase of the network parameters without degrading network performance.

Thus, from the above definition it is clear that the *routing protocol scalability* is dependent on the scalability properties of the network the protocol is run over. That is, the network own scalability properties provides the reference level as to what to expect of a routing protocol. Obviously, if the overhead induced by a routing protocol grows faster than the *network rate* but slower than the *minimum traffic load*, the routing protocol is not degrading network performance, which is being determined by the *minimum traffic load*.

To quantify a *routing protocol scalability*, the respective scalability factor is defined, based on the *total overhead* concept (definition 2), as follows:

Definition 8: Let $X_{ov}(\lambda_1, \lambda_2, \dots)$ be the total overhead induced by routing protocol X , dependent on parameters $\lambda_1, \lambda_2, \dots$ (e.g. network size, mobility rate, data generation rate, etc.). Then, the Protocol X 's routing protocol scalability factor with respect to a parameter λ_i ($\rho_{\lambda_i}^X$) is defined to be :

$$\rho_{\lambda_i}^X \stackrel{\text{def}}{=} \lim_{\lambda_i \rightarrow \infty} \frac{\log X_{ov}(\lambda_1, \lambda_2, \dots)}{\log \lambda_i}$$

The *routing protocol scalability factor* provides a basis for comparison among different routing protocols. Finally, to assess whether a routing protocol is *scalable*, the following definition is used:

Definition 9: A routing protocol X is said to be scalable with respect to the parameter λ_i if and only if, as the parameter λ_i increases, the total overhead induced by such protocol (X_{ov}) does not increase faster than the network's minimum traffic load. That is, if and only if:

$$\rho_{\lambda_i}^X \leq \Psi_{\lambda_i}$$

Thus, for the class of network under study, a routing protocol X is *scalable* with respect to network size if and only if $\rho_N^X \leq 1.5$; it is *scalable* w.r.t. mobility rate if and only if $\rho_{\lambda_{lc}}^X \leq 0$; and it is *scalable* w.r.t. traffic if and only if $\rho_{\lambda_t}^X \leq 1$.

In the remainder of this paper we will derive asymptotic expressions for the total overhead (and therefore the *routing protocol scalability factor*) induced by a representative set of routing protocols. The methodology to be employed consists of computing each of the three components of *total overhead*, namely *proactive*, *reactive* and *sub-optimal routing*, separately and then adding them up. Besides the trivial result that Plain Flood-

⁶This observation is the main reason behind our focusing on networks with power control, where the transmission range is kept in line so that the network degree is kept bounded.

ing (PF) is the only protocol that is *scalable* with respect to mobility, and that most protocols are *scalable* with respect to traffic, the more interesting result that HSLS is *scalable* with respect to network size is found.

III. PLAIN FLOODING (PF)

In PF, each packet is (re)transmitted by every node in the network (except the destination). Thus, $N - 1$ transmissions are required for each data packet, when the optimal value (on average) should have been L . Since there are $\lambda_t N$ data packets generated each second, the additional bandwidth required for transmission of all these packets is *data* $(N - 1 - L)\lambda_t N$ bps. Since $L = \Theta(\sqrt{N})$, the PF's *sub-optimal routing- and total-overhead* per second is equal to $\Theta(\lambda_t(N^2 - N^{1.5})) = \Theta(\lambda_t N^2)$. In consequence $\rho_{\lambda_t}^{PF} = 1$, $\rho_{\lambda_{lc}}^{PF} = 0$, and $\rho_N^{PF} = 2$.

IV. STANDARD LINK STATE (SLS)

In SLS, a node sends a Link State Update (LSU) to the entire network each time it detects a link status change. A node also sends periodic, soft-state LSUs every T_p seconds. There is no *reactive* overhead associated with SLS, and since the paths determined are optimal, there is no *sub-optimal routing* overhead associated with it either.

In SLS, each node generates a LSU at a rate of λ_{lc} per second, so in average there are $N \lambda_{lc}$ LSUs being generated at any given second. Each LSU is retransmitted at least once per each node (i.e. N times), inducing an overhead of *lsu* N bits (where *lsu* is the size of the LSU packet). Then SLS *proactive* and *total-overhead* per second is *lsu* $\lambda_{lc} N^2$ bps, that is, $\Theta(\lambda_{lc} N^2)$; and $\rho_{\lambda_t}^{SLS} = 0$, $\rho_{\lambda_{lc}}^{SLS} = 1$, and $\rho_N^{SLS} = 2$.

V. DYNAMIC SOURCE ROUTING (DSR)

In DSR no proactive information is exchanged. A node (source) reaches a destination by flooding the network with a route request (RREQ) message. When a RREQ message reaches the destination (or a node with a cached route towards the destination) a route reply message is sent back to the source, including the newly found route. The source attaches the new route to the header of all subsequent packets to that destination, and any intermediate node along the route uses this attached information to determine the next hop in the route. The present work focuses on DSR without the route cache option (DSR-noRC). A lower bound for DRS-noRC's *total overhead* is derived next.

The DSR-noRC *reactive* overhead must account for RREQ messages generated by new session requests (at a rate λ_s per second per node) and the RREQ messages generated by failures in links that are part of a path currently in use. If we only consider the RREQ messages generated by new session requests, then a lower bound can be obtained.

Each route request message is flooded to the entire network, resulting in $N - 1$ retransmissions (only the destination does not need to retransmit this message). Thus, each message induces an overhead of *size_of_RREQ* $(N - 1)$ bits, and there are $\lambda_s N$ RREQ messages generated every second due to new session requests. Thus, the DSR-noRC reactive overhead per second is $\Omega(\lambda_s N^2)$.

For the DSR-noRC *sub-optimal routing* overhead a lower bound will be obtained by considering only the extra bandwidth

required for appending the source-route in each data packet. The number of bits appended in each data packet will be proportional to the length L_i of path i . Since this length is not shorter than L_i^{opt} (the optimal path length), using L_i^{opt} instead of L_i will result on a lower bound. The extra bandwidth consumed by a packet delivered using a path i (with at least L_i^{opt} retransmissions) will be at least $(\log_2 N)(L_i^{opt})^2$, where $\log_2 N$ is the minimum length of a node address. The average extra bandwidth per packet over all paths is $E\{(\log_2 N)(L_i^{opt})^2\} \geq (\log_2 N)E\{L_i^{opt}\}^2 = (\log_2 N)L^2$ bits. Thus, for each packet sent from a source to a destination there is an average *sub-optimal routing* overhead of at least $(\log_2 N)L^2$ bits. Since $\lambda_t N$ packets are transmitted per second, the *sub-optimal routing* overhead induced over the entire network is at least $\lambda_t N(\log_2 N)L^2$ bps. Recalling that $L = \Theta(\sqrt{N})$ (assumption a.4), the DSR-noRC *sub-optimal routing* overhead per second is found to be $\Omega(\lambda_t N^2 \log_2 N)$ bps.

Combining the previous results, DSR-noRC *total overhead* per second is $\Omega(\lambda_s N^2 + \lambda_t N^2 \log_2 N)$. Also, $\rho_{\lambda_t}^{DSR-noRC} = 1$, $0 < \rho_{\lambda_{tc}}^{DSR-noRC} \leq 1$,⁷ and $\rho_N^{DSR-noRC} > 2$.

VI. HIERARCHICAL LINK STATE (HIERLS)

In the m -level HierLS routing, network nodes are regarded as level 1 nodes, and level 0 clusters. Level i nodes are grouped into level i clusters, which become level $i+1$ nodes, until the number of highest level nodes is below a threshold and therefore they can be grouped (conceptually) into a single level m . Thus, the value of m is determined dynamically based on the network size, topology, and threshold values.

Link state information inside a level i cluster is aggregated (limiting the rate of LSU generation) and transmitted only to other level i nodes belonging in the same level i cluster (limiting the scope of the LSU). Thus, a node link change may not be sent outside the level 1 cluster (if they do not cause a significant change to higher levels aggregated information), greatly reducing the proactive overhead.

HierLS relies on the Location Management service to inform a source node S of the address of the highest level cluster that contains the desired destination D and does not contain the source node S . For example, consider a 4-level network as shown in Figure 1. S and D are level 1 nodes; $X.1.1$, $X.1.2$, etc. are level 2 nodes (level 1 clusters); $X.1$, $X.2$, etc. are level 3 nodes (level 2 clusters); X , Y , V , and Z are level 4 nodes (level 3 clusters); the entire network forms the level 4 cluster. The Location Management (LM) service provides S with the address of the highest level cluster that contains D and does not contain S (e.g. the level 3 cluster Z in Figure 1). Node S can then construct a route toward the destination. This route will be formed by a set of links in node S level 1 cluster ($X.1.1$), a set of level 2 links in node S level 2 clusters ($X.1$), and so on. In Figure 1 the route found by node S is :

⁷DSR's *total overhead* does depend on mobility, since breakages of links forming existing routes will trigger route discovery procedures that will induce reactive overhead and/or cause route degradation. Similar to the lower bound derived in this section, an upper bound for DSR's *total overhead* may be derived by assuming that each link breakage triggers a global route discovery (regardless of the link being part of an active route or not). Such an upper bound would increase linearly with the mobility rate, and therefore we obtain the upper bound for $\rho_{\lambda_{tc}}^{DSR-noRC} \leq 1$.

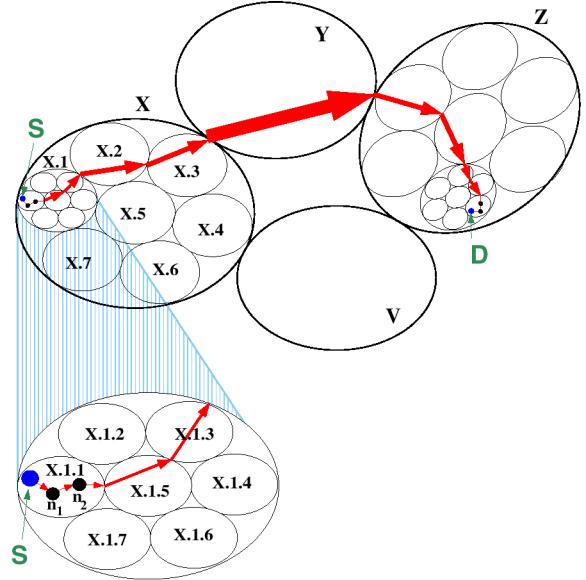


Fig. 1. A Source (S) - Destination (D) path in HierLS.

$S - n_1 - n_2 - X.1.5 - X.1.3 - X.2 - X.3 - Y - Z - D$. When a node outside node S level 1 cluster receives the packet, the node will likely produce the same high-level route towards D , and will 'expand' the high-level links that traverse its cluster using lower level (more detailed) information. In Figure 1 this expansion is shown for the segment $Z - D$. The Location Management (LM) service can be implemented in different ways, whether proactive (location update messages), reactive (paging), or hybrid. Typical choices are:

- LM1: Pure reactive. Whenever a node changes its level i clustering membership but remains in the same level $i+1$ cluster, this node sends an update to all the nodes inside its level $i+1$ cluster. For example, (see Figure 1) if node n_2 moves inside cluster $X.1.5$, i.e. it changes its level 1 cluster membership but does not change its level 2 cluster membership (cluster $X.1$), then node n_2 will send a location update to all the nodes inside cluster $X.1$. The remaining nodes will not be informed.
- LM2: Local paging. In this LM technique, one node in each level 1 cluster assumes the role of a LM server. Also, one node among the level 1 LM servers inside the same level 2 cluster assumes the role of a level 2 LM server, and so on up to level m . The LM servers form a hierarchical tree. Location updates are only generated and transmitted between nodes in this tree (LM servers). When a node D changes its level i clustering membership, the LM server of its new level i cluster will send a location update message to the level $i+1$ LM server, which in turn will forward the update to all the level i LM servers inside this level $i+1$ cluster. Additionally, the level $i+1$ LM server checks if the node D is new in the level $i+1$ cluster, and if this is the case it will send a location update to its level $i+2$ LM server, and so on.

When a level i LM server receives a location update message regarding node D from its level $i+1$ LM server, it updates its local database with node D 's new location information and forwards this information to all the level $i-1$ LM servers inside its level i cluster. Each of these level $i-1$ LM servers forwards the location update message to the level $i-2$ servers in its level

$i - 1$ cluster, and so on until all the level 1 LM servers (inside node D 's level $i + 1$ cluster) are informed of the new level i location information of node D . When a node needs location information about any node in the network, the node pages its level 1 LM server for this information.

- LM3: Global paging. LM3 is similar to LM2. In LM3, however, when a level i LM server receives a location update from a higher level $i + 1$ LM server, it does not forward this information to the lower level ($i - 1$) LM servers. Thus, a lower level (say level $j < i$) LM server does not have location information for nodes outside its level j cluster. A mechanism for removing outdated location information about nodes that left a level i cluster need to be added to the level i clusters LM servers. Basically, a level 1 LM server that detects that a node left its level 1 cluster will remove the entry corresponding to this node from its own database, and will inform its level 2 LM server. The level 2 LM server will wait for a while for a location update from the new level 1 cluster (if inside the same level 2 cluster) and if no such an update is received it will remove the node entry and will inform its level 3 LM server, and so on until arriving to a LM server that already has information about the new location of the node. When a node needs location information about any node in the network, the node pages its level 1 LM server for the information. If the level 1 LM does not have the required information, it (the level 1 LM server) pages its level 2 LM server, who in turn pages its level 3 LM server, and so on, until a LM server with location information about the desired destination is found.

Approach LM1, the easiest to implement, will induce greater overhead and lower latencies for route establishment. Approach LM2 potentially reduces the bandwidth consumption (for reasonable values of λ_s) but at the expense of complexity (selection and maintenance of LM servers) and an increase in the latency associated with route establishment. However, the asymptotic characteristic of HierLS are identical under LM1 and LM2, as will be seen later. Approach LM3 is the more complex to implement. It will induce a significant amount of reactive overhead, but will reduce the amount of overhead induced by mobility. In this paper, results for the HierLS *total overhead* for all three LM Techniques are presented in Table I. However, due to space constraints, only the derivation for the total overhead expression for a HierLS-LM1 (pure proactive LM technique) will be presented next. The reader is referred to [14] or [15] for the remaining derivations.

A. HierLS-LM1 proactive overhead

A network organized in m level clusters, each of equal size k ($N = k^m$) is considered. Note that k is predefined while m increases with N .

Under assumption a.7, HierLS-LM1's *proactive* asymptotic overhead is dominated by the location management function, that induces an overhead that grows at least as fast as $\Theta(sN^{1.5})$ (explained below), where s is the node relative speed. In the other hand, most of the LSUs updates will correspond to level 1 links, and will be propagated inside the level 1 clusters only, thus, LSU packets will induce a proactive overhead that will only grow as fast as $\lambda_{lc} k N$ (this is, of course, a lower bound).

HierLS-LM1 location management overhead expressions,

can be obtained by considering that the time a node takes to change its level $m - 1$ cluster is directly proportional to the diameter of this level $m - 1$ cluster and inversely proportional to the node's relative speed s . Since the level $m - 1$ cluster size is N/k , then the cluster diameter is $\Theta(\sqrt{N/k})$. Under approach LM1, the new location information will have to be forwarded to all the nodes inside the level m cluster (the entire network). Thus, every node will send a location update message to the entire network (N transmissions) each $\Theta(\sqrt{N/k}/s)$ seconds, inducing an overhead of $\Theta(\sqrt{k}s\sqrt{N})$ bits every second. Adding up all nodes contributions, the proactive overhead per second due to level $m - 1$ clusters membership change is $\Theta(\sqrt{ks}N^{1.5})$. Regarding the location updates generated due to level $m - i$ membership change, it can be seen that a level $m - i$ cluster is k^{i-1} times smaller than a level $m - 1$ cluster, and consequently a level $m - i$ cluster's diameter is $k^{\frac{i-1}{2}}$ times smaller than a level $m - 1$ cluster's diameter. Thus, the generation rate of location updates due to level $m - i$ membership changes is $k^{\frac{i-1}{2}}$ times larger than the rate induced by level $m - 1$ changes. Also, since the new location information will have to be transmitted to all the nodes inside the current level $m - i + 1$ cluster, then the number of transmissions required for each packet decreases by a factor of $k^{-(i-1)}$ with respect to the number of transmissions induced by level $m - 1$ changes, which results in a net reduction of $k^{-\frac{i-1}{2}}$. Then, the overhead due to all location updates is :

$$\begin{aligned} Loc_Upd_Cost &= \Theta(\sqrt{ks}N^{1.5})[1 + k^{-\frac{1}{2}} + k^{-1} + \dots] \\ &= \Theta(\sqrt{ks}N^{1.5}) \frac{1}{1 - \sqrt{1/k}} \end{aligned}$$

Thus, the location management overhead is $\Theta(\lambda_{lc}N^{1.5})$ bps (by assumption a.7, λ_{lc} is proportional to s). Combining this value with the lower bound obtained for the LSU-induced overhead ($\Omega(\lambda_{lc}N)$), it is concluded that the HierLS-LM1 *proactive* overhead is $\Theta(\lambda_{lc}N^{1.5})$.

B. HierLS-LM1 sub-optimal routing overhead

To estimate the *sub-optimal routing* overhead, it is assumed that each level i (beginning with level 2) increases the actual route length by a factor f_i (f_i depends on the value of k , the LSU triggering thresholds, and is typically close to 1, for example $f = 1.05$ means a 5% increase in the route length). Thus, if the optimal path length is l , then the actual path length will be $\prod_{i=2}^{i=m} f_i l$. Let f be the geometric average of the set $\{f_i\}$, that is, $f = (\prod_{i=2}^m f_i)^{\frac{1}{m-1}}$. Then, the *sub-optimal routing* overhead induced by a packet transmission is $data[f^{m-1} - 1]l = data[k^{(\log_k f)(m-1)} - 1]l = data[\frac{N^\delta}{k} - 1]l$, where $\delta = \log_k f$. There are $\lambda_t N$ packets generated each second, thus the average *sub-optimal routing* overhead per second is $data(\frac{N^\delta}{k} - 1)L\lambda_t N$. Since L is $\Theta(\sqrt{N})$, we finally get that the HierLS-LM1 *sub-optimal routing* overhead per second is $\Theta(\lambda_t N^{1.5+\delta})$.

C. HierLS-LM1 total overhead

Combining the previous expressions, the HierLS-LM1 *total overhead* is found to be $\Theta(\lambda_{lc}N^{1.5} + \lambda_t N^{1.5+\delta})$. Also,

$\rho_{\lambda_t}^{HierLS-LM1} = 1$, $\rho_{\lambda_{lc}}^{HierLS-LM1} = 1$, and $\rho_N^{HierLS-LM1} = 1.5 + \delta > 1.5$ (HierLS is *almost* scalable w.r.t. network size).

VII. ZONE ROUTING PROTOCOL (ZRP)

ZRP is a hybrid approach, combining a proactive and a reactive part, trying to minimize the sum of their respective overheads. In ZRP, a node disseminates event-driven LSUs to its k -hop neighbors (nodes at a distance, in hops, of k or less). Thus, each node has full knowledge of its k -hop neighborhood and may forward packets to any node within it. When a node needs to forward a packet outside its k -hop neighborhood, it sends a route request to a subset of the nodes in the network, namely the ‘border nodes’. The ‘border nodes’ will have enough information about their k -hop neighborhoods to decide whether to reply to the route request or to forward it to its own set of ‘border’ nodes. The route formed will be described in terms of the ‘border’ nodes only, thus allowing ‘border’ nodes to locally recover from individual link failures, reducing the overhead induced by route maintenance procedures.

The following lower bound for ZRP’s *total overhead* (ZRP_{ov}) was obtained:

$$ZRP_{ov} = \begin{cases} \Omega(\lambda_{lc} N^2) & \text{if } \lambda_{lc} = O(\lambda_s / \sqrt{N}) \\ \Omega(\lambda_{lc}^{\frac{1}{3}} \lambda_s^{\frac{2}{3}} N^{\frac{5}{3}}) & \text{if } \lambda_{lc} = \Omega(\lambda_s / \sqrt{N}) \text{ and } \lambda_{lc} = O(\lambda_s N) \\ \Omega(\lambda_s N^2) & \text{if } \lambda_{lc} = \Omega(\lambda_s N) \end{cases}$$

Due to space limitations, the derivation of the ZRP *total overhead* was left out of the paper. Once again, the reader is referred to [14] or [15] for the complete derivation.

Note that the asymptotic expression provides us with much more information about the parameters interactions than the scalability factors, which are computed assuming that just one parameter is increased while the others remain fixed. For ZRP, $\rho_{\lambda_t}^{ZRP} = 0$ (pure proactive mode), $0 < \rho_{\lambda_{lc}}^{ZRP} \leq 1$ (pure reactive mode, similar to DSR’s), and $\rho_N^{ZRP} \geq 1.66$. Note that the information provided by the *scalability factors* is incomplete, and it hinds the fact that the exponential rates of increase of ZRP’s total overhead with respect to mobility and traffic always add up to at least 1, as can be seen from the *total overhead*’s asymptotic expressions.

VIII. HAZY SIGHTED LINK STATE (HSLS)

HSLS is based on the observation that nodes that are far away do not need to have complete topological information in order to make a good next hop decision. Thus, propagating every link status change over the entire network may not be necessary. In a highly mobile environment, a node running HSLS will transmit - provided that there is a need to - a LSU only at particular time instants that are multiples of t_e seconds. Thus, potentially several link changes are ‘collected’ and transmitted every t_e seconds. The *Time To Live* (TTL) field of the LSU packet is set to a value (which specifies how far the LSU will be propagated) that is a function of the current time index as explained below. After one global LSU transmission – LSU that travels over the entire network, i.e. TTL field set to infinity, as for example during initialization – a node ‘wakes up’ every t_e seconds and sends a LSU with TTL set to 2 if there has been a link status change in

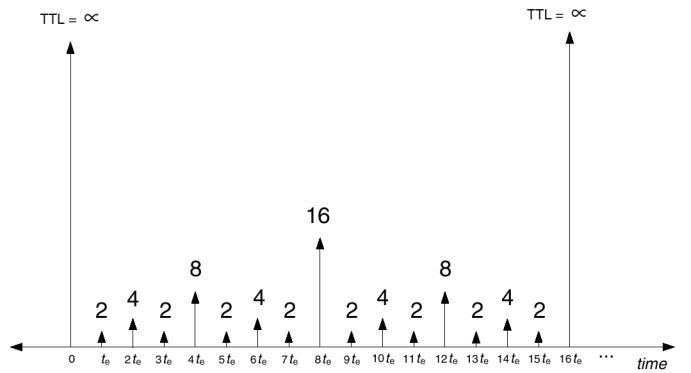


Fig. 2. HSLS’s LSU generation process (mobility is high).

the last t_e seconds. Also, the node wakes up every $2t_e$ seconds and transmits a LSU with TTL set to 4 if there has been a link status change in the last $2t_e$ seconds. In general, a node wakes up every $2^{i-1}t_e$ ($i = 1, 2, 3, \dots$) seconds and transmits a LSU with TTL set to 2^i if there has been a link status change in the last $2^{i-1}t_e$ seconds. If a packet TTL field value (2^i) is greater than the distance from this node to any other node in the network (which will cause the LSU to reach the entire network), the TTL field of the LSU is reset to infinity (global LSU), and the algorithm is re-initiated.

Nodes that are at most two hops away from a node, say X , will receive information about node X ’s link status change at most after t_e seconds. Nodes that are more than 2 but at most 4 hops away from X will receive information about any of X links change at most after $2t_e$ seconds. In general, nodes that are more than 2^{i-1} but at most 2^i hops away from X will receive information about any of X links change at most after $2^{i-1}t_e$ seconds. Figure 2 shows an example of HSLS’s LSU generation process when mobility is high and in consequence LSUs are always generated. An arrow with a number over it indicates that at that time instant a LSU (with TTL field set to the indicated value) was generated and transmitted. Figure 2 assumes that the node executing HSLS computes its distance to the node farthest away to be between 17 and 32 hops, and therefore it replaces the TTL value of 32 with the value infinity, resetting the algorithm at time $16t_e$. The reader is referred to [11] and [12] for more details about HSLS.

A. HSLS proactive overhead

A highly mobile environment (i.e. a LSU is generated every time interval) is considered. All the different LSUs (re)transmissions due to LSUs generated by a node, say X , will be added and then averaged over time. The value obtained will be multiplied by the number of nodes in the network to get the *proactive overhead*. LSUs will be grouped based on their TTL value at the time they were generated, beginning with the LSUs with larger TTL values.

Let MD_x be the maximum distance from node X to any other node in the network. Let R_x be the power of 2 such that $R_x < MD_x \leq 2R_x$. For example, $R_x = 16$ in figure 2, where MD_x was assumed to be between 17 and 32. Under HSLS, node X computes MD_x each t_e seconds based on its own topology information, which is not necessarily up-to-date, so MD_x

is a time-changing value that is not being timely updated. The above observation, however, will have little impact on the value of R_x , which may be assumed roughly constant over time.

Let's consider what happens at time $R_x t_e$ ($16t_e$ in figure 2). At this time node X sends a LSU to the entire network and the algorithm is re-initiated. Thus, every $R_x t_e$ seconds node X induces N transmissions, and therefore the bandwidth consumption due to these global LSUs is $\frac{lsu N}{R_x t_e}$, where lsu is the average length of a LSU packet.

The second larger TTL is R_x , and LSUs with this TTL are generated $\frac{R_x}{2} t_e$ seconds after a global LSU is sent (times $8t_e$ in figure 2). Recalling that the timers are reset at time $R_x t_e$, we notice that the interval between consecutive generation times is $(R_x t_e - \frac{R_x}{2} t_e) + \frac{R_x}{2} t_e = R_x t_e$. Thus, the generation rate of LSUs with TTL equal to R_x is $\frac{1}{R_x t_e}$ (the same as the generation rate of global LSUs). These LSUs will not reach all the nodes in the network but only a fraction f_x . From assumption a.3, f_x should be around $(R_x/Md_x)^2$, i.e., $f_x \in [0.25, 1]$. In practical situations, due to boundary effects (i.e. the number of nodes at a maximum distance Md_x is small), we obtain that typically f_x is in the interval $[0.5, 1]$. Thus, the bandwidth consumption due to LSUs with TTL equal to R_x is $\frac{lsu f_x N}{R_x t_e}$.

For the remaining TTL values, ‘boundary’ conditions are no longer relevant. Thus, for TTL equal to $R_x/2$ the generation rate doubles (e.g. LSUs with TTL equal to 8 are sent at times $4t_e, 12t_e, \dots$ in figure 2), and the number of transmissions induced per LSU is reduced by a factor of 4 (because of assumption a.3, and the fact that the TTL values are reduced to a half); thus the total effect is a reduction by a factor of 2 with respect the bandwidth consumption due to LSUs with TTL equal to R_x . The same argument applies for TTL equal to $R_x/4, R_x/8, \dots, 2, 1$.⁸ Finally, the total bandwidth consumption due to all the LSUs generated by node X is equal to :

$$\begin{aligned} X_{HSLS}^{pro} &= \frac{lsu N}{R_x t_e} + \frac{lsu f_x N}{R_x t_e} + \frac{lsu f_x N}{2R_x t_e} + \frac{lsu f_x N}{4R_x t_e} + \dots \\ &= \frac{lsu N}{R_x t_e} [1 + f_x(1 + \frac{1}{2} + \frac{1}{4} + \dots)] \approx \frac{lsu N}{R_x t_e} [1 + 2f_x] \end{aligned}$$

Since the size of a LSU depends only on the node density (bounded on average), f_x is bounded below 1, and R_x is $\Theta(\sqrt{N})$ (assumption a.4); the proactive overhead per second induced by one node is $\Theta(\frac{N^{0.5}}{t_e})$. Since there are N nodes, the proactive overhead per second induced by the entire network is $\Theta(\frac{N^{1.5}}{t_e})$.

B. HSLS sub-optimal routing overhead

Due to space constraints, the complete derivation was left out of the paper. Below, an insight into it is provided. The reader is referred to [14] or [15] for the actual derivation.

Let t_k^{elap} be the maximum time elapsed since ‘fresh’ LSU information about a destination k hops away was last received. HSLS induces a quasi-linear relationship between t_k^{elap} and k . In general, $\frac{t_e}{2} \leq \frac{t_k^{elap}}{k} \leq t_e$. Thus, the ratio between the time

⁸ Assumptions a.3 and a.4 are asymptotic conditions, and as such, are not applicable to small values of TTL. However, the contributions of LSUs with small TTL values in the proactive overhead of a large network is not significant and a more exact analysis can be safely omitted.

elapsed since fresh information was received and distance is bounded by t_e , independently of network size or distance to the destination. Based on the mobility model assumption a.8 (time scaling), this will cause the probability of a sub-optimal next hop decision to be bounded⁹, and the fraction of the increase of the sub-optimal routes (with respect to the optimal ones) to also be bounded independently of network size. Then, for a fixed value of t_e , HSLS sub-optimal routing overhead will increase as $\Theta(\lambda_t N^{1.5})$.

To investigate the dependence of the sub-optimal routing overhead on the time t_e , a more precise mobility model need to be defined. Assuming a mobility model that induces an exponential residence time on a given area, HSLS sub-optimal routing overhead was found to be equal to : $\Theta((e^{\lambda_{lc} t_e K_4} - 1)\lambda_t N^{1.5})$, where K_4 is a constant.

C. HSLS total overhead

There is no reactive overhead associated with HSLS. Thus, the HSLS total overhead for the class of networks analyzed in the previous subsections is equal to :

$$HSLS_{ov} = N^{1.5} [K_5 \frac{1}{t_e} + K_3 (e^{\lambda_{lc} t_e K_4} - 1)\lambda_t]$$

The value of t_e should be tuned to optimize performance. For a moment, let's use the approximation $e^x - 1 \approx x$, where $x = \lambda_{lc} t_e K_4$. Thus:

$$HSLS_{ov} \approx N^{1.5} [\frac{K_5}{t_e} + K_6 \lambda_{lc} \lambda_t t_e]$$

Choosing the value of t_e that minimizes the above expression we get $t_e = \Theta(\frac{1}{\sqrt{\lambda_{lc} \lambda_t}})$, $x = \Theta(\frac{\sqrt{\lambda_{lc}}}{\sqrt{\lambda_t}})$, and $HSLS_{ov} = \Theta(\sqrt{\lambda_{lc} \lambda_t} N^{1.5})$. The previous expression would define the asymptotic behavior of HSLS’s total overhead only if our approximation $e^x - 1 \approx x$ is valid. Indeed, if λ_t grows asymptotically faster than λ_{lc} , the value of x goes to zero and the approximation $e^x - 1 \approx x$ is valid. On the other hand, if λ_{lc} grows asymptotically faster than λ_t , the approximation will not be valid. In this case, since the exponential function is the fastest growing, it is desirable to maintain the product $\lambda_{lc} t_e$ (and therefore the value of p) bounded and therefore we choose $t_e = \Theta(\frac{1}{\lambda_{lc}})$. Thus, the HSLS total overhead in this scenario becomes $\Theta(N^{1.5}(\lambda_{lc} + \lambda_t)) = \Theta(\lambda_{lc} N^{1.5})$, where the last equality holds due to our assumption that λ_{lc} grows asymptotically faster than λ_t and therefore λ_{lc} dominates the previous expression. Thus, the HSLS’s total overhead is :

$$HSLS_{ov} = \begin{cases} \Theta(\sqrt{\lambda_{lc} \lambda_t} N^{1.5}) & \text{if } \lambda_{lc} = O(\lambda_t) \\ \Theta(\lambda_{lc} N^{1.5}) & \text{if } \lambda_{lc} = \Omega(\lambda_t) \end{cases}$$

Also, it can be noted that $\rho_{\lambda_t}^{HSLS} = 0.5$, $\rho_{\lambda_{lc}}^{HSLS} = 1$, and $\rho_N^{HSLS} = 1.5$. Thus, HSLS is the only protocol that is scalable with respect to network size.

⁹ Since the ratio maximum displacement – speed times elapsed time – over distance is bounded, so is the ‘angular’ displacement of the destination. The ‘angular’ displacement will determine whether the node chosen as the next hop is the proper one or not.

IX. COMPARATIVE STUDY

In the previous sections the *scalability factors* of several representative routing protocols have been derived. From those results we concluded that PF is the only protocol known to be *scalable* w.r.t. mobility ($\rho_{\lambda_{tc}}^{PF} = 0$), while all of the protocols were *scalable* w.r.t. traffic. More interesting was to find that HSLS is the only protocol *scalable* with respect to network size ($\rho_N^{HSLS} = 1.5$). However, much more information about the protocol parameter's interactions may be derived from the asymptotic *total overhead* expressions, which are summarized in Table I.

Table I presents our results for the *total overhead* when the tunable parameters are selected to optimize performance (or at least, optimize the lower bounds derived before). These results increase our understanding of the limits and provide valuable insight about the behavior of several representative routing protocols. The better understanding of these limits will help network designers to better identify the class of protocols to engage depending on their operating scenario. For example, if the designer's main concern is network size, it can be noted that HierLS and HSLS scale better than the others. Similarly, if traffic intensity is the most demanding requirement, then SLS and ZRP are to be preferred since they scale better with respect to traffic (*total overhead* is independent of λ_t); HSLS follows as it scales as $\Theta(\sqrt{\lambda_t})$, and PF, DSR, and HierLS are the last since their *total overhead* increases linearly with traffic.¹⁰

Similarly with respect to the rate of topological change, we observe that PF may be preferred (if size and traffic are small and the rate of topological change increases too rapidly), since its *total overhead* is independent of the rate of topological change. Provably next will be ZRP and DSR since their lower bounds are independent of the rate of topological changes. The bounds are not necessarily tight, and ZRP's and DSR's behavior should depend somewhat of the rate of topological change. Finally, for SLS, HierLS, and HSLS we know (as opposed to DSR and ZRP where we suppose) that their *total overhead* increase linearly with the rate of topological change.

It is interesting to note that when only the traffic or the mobility is increased (but not both), ZRP can achieve almost the best performance in each case.¹¹ However, if mobility and traffic increase at the same rate; that is, $\lambda_{tc} = \Theta(\lambda)$ and $\lambda_t = \Theta(\lambda)$ (for some parameter λ), then ZRP's *total overhead* ($\Omega(\lambda N^{1.66})$) will present the same scalability properties as HSLS's ($\Theta(\lambda N^{1.5})$) and HierLS's ($\Theta(\lambda N^{1.5+\delta})$) with respect to λ , with the difference that ZRP does not scale as well as the other two with respect to size.

These and more complex analyses can be derived from the expression presented in this paper, when different parameters are modified simultaneously accordingly with the scenario the designer is interested in.

¹⁰It is interesting to note that HSLS scales better with traffic intensities than HierLS (the only other protocol that scales well with size). This result may have an intuitive explanation in the fact that HierLS never attempts to find optimal routes towards the destination, even under slowly changing conditions. HSLS on the other hand, may eventually obtain full topology information – and therefore optimal routes – if the rate of topological changes is small with respect to $1/t_e$, as is the case when λ_t grows faster than λ_{tc} .

¹¹Almost, because ZRP can not achieve the independence of *total overhead* from mobility. PF does.

Proto.	Total over. (best)	Cases
PF	$\Theta(\lambda_t N^2)$	Always
SLS	$\Theta(\lambda_{tc} N^2)$	Always
DSR	$\Omega(\lambda_s N^2 + \lambda_t N^2 \log_2 N)$	no Route Cache
HierLS	$\Theta(\lambda_{tc} N^{1.5} + \lambda_t N^{1.5+\delta})$ $\Theta(\lambda_{tc} N \log N + \lambda_t N^{1.5+\delta})$	LM1 or LM2 LM3
ZRP	$\Omega(\lambda_{tc} N^2)$ $\Omega(\lambda_s N^2)$ $\Omega(\lambda_{tc}^{\frac{1}{3}} \lambda_s^{\frac{2}{3}} N^{\frac{5}{3}})$	$\lambda_{tc} = O(\lambda_s / \sqrt{N})$ $\lambda_{tc} = \Omega(\lambda_s N)$ otherwise
HSLS	$\Theta(\sqrt{\lambda_{tc} \lambda_t} N^{1.5})$ $\Theta(\lambda_{tc} N^{1.5})$	$\lambda_{tc} = O(\lambda_t)$ $\lambda_{tc} = \Omega(\lambda_t)$

TABLE I
ASYMPTOTIC TOTAL OVERHEAD EXPRESSIONS.

HSLS has better asymptotic properties than HierLS, which means that as size increases HSLS eventually outperform HierLS. The idea of HSLS – being much more simple to implement – outperforming HierLS is counter-intuitive. A first reaction to this result will likely be to assume that the constants involved in the asymptotic analysis may be too large, preventing HSLS from outperform HierLS under ‘reasonable’ scenario. Thus, the authors relied on a couple of simulation experiment to validate if, in effect, HSLS may outperform HierLS even under moderate network size and traffic load.

A. A simulation experiment: HSLS vs. HierLS-LM1

Table II shows the simulation results obtained by OPNET for a 400-node network where nodes are randomly located on a square of area equal to 320 square miles (i.e. density is 1.25 nodes per square mile). Each node choose a random direction among 4 possible values, and move on that direction at 28.8 mph. Upon reaching the area boundaries, a node bounces back. The radio link capacity was 1.676 Mbps. Simulation were run for 350 seconds, leaving the first 50 seconds for protocol initialization, and transmitting packets (60 8kbps streams) for the remaining 300 seconds. The HierLS approach simulated was of the type HierLS-LM1, and corresponds to the DAWN project [10] modification of the MMWN clustering protocol [8]. The minimum and maximum cluster size were set to 9 and 35 respectively.

The metric of interest is the throughput (i.e. fraction of packets successfully delivered). The simulation results presented are not a comprehensive study of the relative performance of HierLS versus HSLS under all possible scenarios. They just presents and example of a real-life situation where HSLS outperform HierLS, and complement our theoretical analysis. The theoretical analysis focuses on asymptotically large network, heavy traffic load, and saturation conditions where the remaining capacity determines the protocol performance. The simulation results, in the other hand, refer to medium size networks with moderate loads, where depending on the MAC employed, other factors may have more weight over the protocols performance.

Table II shows the throughput obtained under two different MAC protocols: unreliable and reliable CSMA. For reliable

Protocol	UNRELIABLE	RELIABLE
HSLS	0.2454	0.7991
HierLS-LM1	0.0668	0.3445

TABLE II
THROUGHPUT OF A 400-NODE NETWORK.

CSMA, packets were retransmitted up to 10 times if a MAC-level ACK was not received in a reasonable time. We can see that in both cases HSLS outperforms HierLS, although the relative difference is reduced under the reliable MAC case. This can be explained considering that the high rate of collisions experienced under unreliable CSMA favored shorter paths. For nodes close by, HSLS may provide almost optimal routes while HierLS routes may be far from optimal if the destination belong to a neighboring cluster. Thus, we can see that unreliable MAC biases towards HSLS. Another factor to take into account is the latency to detect link up/downs. Under HierLS this information is synchronized among all the nodes in the cluster and therefore some latency is enforced to avoid flapping. In HSLS, in the other hand, each node may have its own view of the network, and as a consequence a node may be more aggressive in temporarily turning links down without informing other nodes. As a consequence, HSLS is more aggressive and reacts much faster to link degradation, using alternate paths if available.

It can be seen that the previous results are highly influenced for another factors such as the MAC protocol being used, the quality of the links that neighbor discovery declares up, the latency on detecting link failures, etc. So, whether HSLS or HierLS should be preferred for a particular scenario, depends on the particular constraints (for example, if memory or processing time is an issue, HierLS may be preferred since it require to store/process a smaller topology table). The present work, however, provides some guidelines, suggesting that as traffic, network size, and data rate increases, and a better MAC is employed (allowing to achieve the full channel capacity), HSLS should tend to be preferred.

X. CONCLUSIONS

The applications for ad hoc networking are only beginning to be recognized. However, before practical implementations are possible, it is necessary to design scalable systems. Hence, scalability has become a dominant objective of ad hoc network algorithm designers. Unfortunately, the community lacks a basic tenet for understanding the fundamental limitations and invariants associated with ad hoc networks.

This paper addresses this shortcoming by presenting a novel and powerful framework (the *total overhead* criteria) that allows for an analytical comparison, and deeper understanding of the characteristics and tradeoffs associated with various classes of routing protocols for mobile networks. This framework, first introduced in [11] to analyze a family of link state protocol variants, was fully developed in this paper and applied to study a variety of protocols that have been proposed and analyzed via simulation methodologies in the literature.

The analytical methods developed in this paper and the result-

ing asymptotic analysis of *total overhead* provide an important contribution to the field that promises to shed new light on the fundamental limitations and underlying characteristics of mobile networks in general, and in the studied protocols in particular. It was found that, among the protocol studied, PF is the only protocol that scales w.r.t. mobility, all of them scale w.r.t. traffic, and HSLS is the only one that scales w.r.t. network size (note that HierLS *almost* scale w.r.t. network size). Thus, the results for HSLS – a novel, easy-to-implement link state variant – showed that the implementation of a complex hierarchy was not mandatory for scalability. A more focused comparison between HierLS and HSLS was undertaken, and as a result, HSLS was established as a competitive alternative to HierLS.

Finally, this work is only a first step. Greater understanding is required of cross-layer interactions and the impact of more general mobility models and traffic workloads. We hope this work will help to lay a foundation for a renewed approach to research into ad hoc networks. The success of this technology depends on rigorous techniques and proof of concepts.

REFERENCES

- [1] J. Broch, D. Maltz, D. Johnson, Y. Hu, and J. Jetcheva, "A Performance Comparison of Multihop Wireless Ad Hoc Network Routing Protocols," *Proceedings of MOBICOM'98*, Dallas, TX., October 1998.
- [2] V. D. Park, and S. Corson, "A Performance Comparison of the Temporally-Ordered Routing Algorithm and Ideal Link-State Routing," *Proceedings of IEEE Symposium on Computers and Communications ISCC'98*, Athens, Greece, June 1998.
- [3] C. E. Perkins, E. M. Royer, S. R. Das, and M. K. Marina, "Performance Comparison of Two On-Demand Routing Protocols for Ad Hoc Networks", *IEEE Personal Communications Magazine*, Vol. 8, No. 1, Feb. 2001.
- [4] P. Jacquet and L. Viennot, 'Overhead in Mobile Ad-hoc Network Protocols", *INRIA Research Report 3965*, Institut National de Recherche en Informatique et en Automatique (INRIA), France, June 2000.
- [5] P. Gupta and P.R. Kumar. "The Capacity of Wireless Networks", *IEEE Transaction on Information Theory*, 46 (2):388-404, March 2000.
- [6] M. Grossglauser and D. Tse. "Mobility Increases the Capacity of Ad-hoc Wireless Networks", in *Proceedings of IEEE Infocom'2001*, Anchorage, Alaska, April 2001.
- [7] D. B. Johnson and D. Maltz, "Dynamic Source Routing in Ad Hoc Wireless Networks.", In Mobile Computing, edited by Tomasz Imielinski and Hank Korth. Kluwer Academic Publishers, 1995.
- [8] S. Ramanathan, M. Steenstrup, "Hierarchically-organized, Multihop Mobile Networks for Multimedia Support", *ACM/Baltzer Mobile Networks and Applications*, Vol. 3, No. 1, pp 101-119.
- [9] Z. Haas and M. Pearlman, "The performance of query control schemes for the zone routing protocol," in *ACM SIGCOMM*, 1998.
- [10] <http://www.ir.bbn.com/projects/dawn/dawn-index.html>
- [11] C. Santivanez, S. Ramanathan, and I. Stavrakakis, "Making Link State Routing Scale for Ad Hoc Networks", In *Proceedings of MobiHOC'2001*, Long Beach, CA, Oct. 2001.
- [12] C. Santivanez, R. Ramanathan, "Hazy Sighted Link State (HSLS) Routing: A Scalable Link State Algorithm", *BBN technical memo BBN-TM-1301*, BBN technologies, Cambridge, MA, August 2001. Available at <http://www.ir.bbn.com/documents/techmemos/index.html>
- [13] R. Guerin, et. al., "Equivalent Capacity and Its Applications to Bandwidth Allocation in High Speed Networks," *IEEE Journal of Selected Areas on Communications*, vol. 9, no. 7, pp. 968-981, Sept. 1991.
- [14] C. Santivanez, "Asymptotic Behavior of Mobile Ad Hoc Routing Protocols with respect to Traffic, Mobility, and Size," *Technical Report TR-CDSP-00-52*, Center for Communications and Digital Signal Processing (CDSP), Northeastern University, Boston, MA, October 2000. Available at <http://www.cdsp.neu.edu/info/students/cesar/analysis.ps.gz>
- [15] C. Santivanez, "A framework for multi-mode routing in wireless ad hoc networks: theoretical and practical aspects of scalability and dynamic adaptation to varying network size, traffic and mobility patterns," *Doctoral thesis*, Electrical and Computing Engineering Department, Northeastern University, Boston, MA, November 2001.